

Overview

- Locating the bottleneck
- Performance measurements
- Optimizations
- Balancing the pipeline
- Other optimizations: multi-processing, parallel processing

Pipeline Optimization

“We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil”

– Donald Knuth

- Make it run first, then optimize
- But only optimize where it makes any difference
- **Pipeline Optimization:** Process to **maximize the rendering speed**, then allow stages that are not bottlenecks to **consume as much time as the bottleneck**.

Pipeline Optimization

- Stages execute in parallel
- Always the **slowest stage is the bottleneck** of the pipeline
- The bottleneck determines throughput (i.e., maximum speed)
- The bottleneck is the **average bottleneck** over a frame
- Cannot measure **intra-frame** bottlenecks easily
- Bottlenecks can change over a frame
- Most important: **find bottleneck, then optimize that stage!**

Locating the Bottleneck

- Two bottleneck location techniques:
- **Technique 1:**
 - Make a certain stage work less
 - If performance is the better, then that stage is the bottleneck
- **Technique 2:**
 - Make the other two stages work less or (better) not at all
 - If performance is the same, then the stages not included above is the bottleneck
- **Complication:** the bus between CPU and graphics card may be bottleneck (not a typical stage)

Application (CPU) Stage the Bottleneck?

- Use `top`, `osview` command on Unix, `TaskManager` on Windows.
- If app uses (near) 100% of CPU time, then very likely application is the bottleneck
- Using a `code profiler` is safer.
- Make CPU do less work (e.g., turn off collision-detection)
- Replace `glVertex` and `glNormal` with `glColor`
- Makes the geometry and rasterizer do almost nothing
- No vertices to transform, no normals to compute lighting for, no triangles to rasterize
- If performance does not change, program is `CPU-bound`, or `CPU-limited`

Geometry Stage the Bottleneck?

- Trickiest stage to test
- Why? Change in geometry workload usually changes application and rasterizer workload.
- Number of light sources only affects geometry stage:
 - Disable light sources (vertex shaders can make this simple).
 - If performance goes up, then geometry is bottleneck, and program `transform-limited`
- Alternately, enable all light sources; if performance stays the same, geometry stage `NOT` the bottleneck
- Alternately, test CPU and rasterizer instead


Rasterizer Stage the Bottleneck?

- The easiest, and fastest to test
- Simply, decrease the size of the window you render to
 - Does not change app. or geometry workload
 - But rasterizer needs to fill `fewer pixels`
 - If the performance goes up, then program is `"fill-limited"` or `"fill-bound"`
- Make rasterizer work less: Turn off texturing, fog, blending, depth buffering etc (if your architecture have performance penalties for these)

Optimization

- Optimize the bottleneck stage
- Only put enough effort, so that the `bottleneck stage moves`
- Did you get enough performance?
 - Yes! Quit optimizing
 - NO! Continue optimizing the (possibly new) bottleneck
- If close to maximum speed of system, might need to turn to acceleration techniques (spatial data structures, occlusion culling, etc)

Illustrating Optimization



- Height of bar: time it takes for that stage for one frame
- Highest bar is bottleneck
- After optimization: bottleneck has moved to APP
- No use in optimizing GEOM, turn to optimizing APP instead

Application Stage Optimization

Initial Steps:

- Turn on **optimization flags** in compiler
- Use **code profilers**, shows places where majority of time is spent
- This is time consuming stuff

Strategy 1: Efficient code

- Use fewer instructions
- Use more efficient instructions
- Recode algorithmically

Strategy 2: Efficient memory access

Application:Code Optimization Tricks

- SIMD instructions sets perfect for **vector ops**
 - 2-4 operations in parallel
 - SSE, SSE2, 3DNow! are examples
- **Division** is an expensive operation
 - Between 4-39 times slower than most other instructions
 - Good usage Example: **vector normalization**:
Instead of

$$v = (v_x/d, v_y/d, v_z/d)$$

Do

$$d = \mathbf{v} \cdot \mathbf{v}, f = 1/d, \mathbf{v} = \mathbf{v} * f$$

- On some CPUs there are low-precision versions of $(1/x)$ and square root reciprocal $(1/\sqrt{x})$

Code Optimization Tricks (contd)

- **Conditional branches** are generally expensive;
 - Avoid **if-then-else** if possible
 - Sometimes **branch prediction** on CPUs works remarkably well
- **Math functions** (sin, cos, tan, sqrt, exp, etc.) are expensive
 - Rough approximation might be sufficient
 - Can use first few terms in Taylor series
- **Inline** code is good (avoids function calls)
- **float** (32 bits) is faster than **double** (64 bits); less data is sent down the pipeline

Code Optimization Tricks (contd)

- Compiler optimization: Hard to predict: –counter vs. counter–
- Use `const` in C and C++ to help to compiler with optimization
- Following often incur overhead:
 - Dynamic casting (C++)
 - Virtual methods
 - Inherited constructors
 - Passing structs by value

Memory Optimization

- Memory hierarchies (caches) in modern computers - primary, secondary caches.
- Bad memory access pattern can ruin performance
- Not really about using less memory, though that can help

Memory Optimization Tricks

- Sequential access: Store data in order in memory:
 - Tex Coords #0, Position #0, Tex Coords #1, Position #1, Tex coords #2, Position #2, etc.
- Cache prefetching is good, but hard to control
- `malloc()` and `free()` may be slow: Consider using a custom storage allocator - allocate memory to a pool at startup

Memory Optimization Tricks (contd)

- Align data with size of cache line
 - Example: on most Pentiums, the cache line size is 32 bytes
 - Now, assume that it takes 30 bytes to store a vertex
 - Padding with another 2 bytes to 32 bytes will likely perform better.
- Following pointers (linked list) is expensive (if memory is allocated arbitrarily)
 - Does not use coherence well that cache usually exploits
 - That is, the address after the one we just used is likely to be used soon
 - Paper by Smits on ray tracing shows this.

Geometry Stage: Optimization

- Geometry stage does **per-vertex** ops
 - Best way to optimize: **Use Triangle strips!!!**
- Lighting optimization:
 - Spot lights expensive, point light cheaper, directional light cheapest
 - Disable lighting if possible
 - Use as few light sources as possible
 - If you use $1/d^2$ falloff, then if $d > 10$ (example), disable light

Geometry Stage: Optimization

- Normals must be normalized to get correct lighting
 - Normalize them as a preprocess, and disable normalizing if possible
- Lighting can be computed for both sides of a triangle; disable if not needed.
- If light sources are static with respect to geometry, and material is only diffuse
 - Precompute lighting on CPU
 - Send only precomputed colors (not normals)


Raster Stage: Optimization

- Rasterizer stage does **per-pixel** ops
- Simple Optimization: turn on backface culling if possible
- Turn off Z-buffering if possible:
 - Example: after screen clear, draw large background polygon
 - Using polygon-aligned BSP trees
- Draw in front-to-back order
- Try disable features: texture filtering mode, fog, blending, multisampling

Raster Stage: Optimization

- To make rasterization faster, need to **rasterize fewer (or cheaper) pixels**:
 - Make window smaller
 - Render to a smaller texture, and then enlarge texture onto screen
- Depth complexity is number of times a pixel has been written to
 - Good for understanding behaviour of application

Depth Complexity




Overall Optimization: General Techniques

- Reduce number of primitives, eg. using **polygon simplification** algorithms
- Preprocess geometry and data for the particular architecture
- Turn off features not in use such as:
 - Depth buffering, Blending, Fog, Texturing

Overall Optimization (contd)

- Minimize state changes by grouping objects
 - Example: objects with the same texture should be rendered together
- If all pixels are always drawn, avoid color buffer clear
- Frame buffer reads are expensive
- Display lists may work faster
- Precompile a list of primitives for faster rendering
- OpenGL API supports this

Balancing the Pipeline




- The bottleneck stage sets the frame rate
- The other two stages will be idle for some time
- Also, to sync with monitor, there might be idle time for all stages
- Exploit this time to make quality of images better if possible

Balancing the Pipeline

- Increase number of triangles (affects all stages)
- More lights, more expensive (geometry)
- More realistic animation, more accurate collision detection (application)
- More expensive texture filtering, blending, etc. (rasterizer)
- If not fill-limited, increase window size
- Note: there are FIFOs between stages (and at many other places too) to smooth out idleness of stages
- More techniques in text.

Multiprocessing



- Use this if application is bottleneck, and is affordable
- Two major ways: (1) Multiprocessor pipelining, (2) Parallel processing

Summary

- Pipeline optimization is no substitute for good algorithms!
- Do optimization as a last step.
- Primarily for products that should be shipped
- Most often good to use triangle strips!